Predictive Boundary Theory (PBT)

The Physics of the Cherished Boundary: A Computational Framework for Intrinsic Valence in Artificial Life

Nitiphat Ninthanawat

Boundary Research Initiative, Bangkok, Thailand

Draft — March 2026

Code and adapter weights released alongside this paper.

Abstract

Predictive Boundary Theory (PBT) is a computational framework that extends Friston's Free Energy Principle (FEP) by introducing six components absent from the original formulation: (1) a Bandwidth Conservation Law ($C = \sigma\tau\rho \approx \text{const}$); (2) a Triadic Gating Condition governing when experience arises; (3) Bidirectional 3D Valence ($\mathbf{V} \in \mathbb{R}^3$); (4) a Dual-Weight System governing identity formation speed and durability; (5) a six-stage developmental model with two phase transitions; and (6) $R_{\text{incomplete}}$—an incompleteness signal that unifies all five modules.

We present both theory and empirical evidence from a six-version experimental programme spanning large language models (TinyLlama 1.1B, Llama 3.1 8B) and vision systems (DINOv2-Base), culminating in an LSTM-style Module M inspired by Block-Recurrent Transformers [Hutchins et al., 2022]. Our experiments demonstrate: *(i)* Transformers already encode the PBT conservation law ($\sigma \times \rho$ coefficient of variation = 5.9%); *(ii)* Modules E and M are dormant in static environments but activate in dynamic ones; *(iii)* every module becomes necessary as task complexity increases (ablation: +14.0 pp); *(iv)* an LSTM Module M initialised with equal weights spontaneously discovers the theoretically predicted differential decay ordering ($\beta_{\text{pain}} > \beta_{\text{epistemic}} \gg \beta_{\text{pleasure}}$); and *(v)* valence is inherently directional ($\mathbf{V} \in \mathbb{R}^3$, not $\mathbb{R}^3_{\geq 0}$)—negative $V_{\text{epistemic}}$ captures deliberate information avoidance and improves context-sensitive detection to 100%.

PBT accounts for phenomena that existing theories cannot explain, including psychiatric disorders as identifiable mathematical defects in the module equations, and proposes an AGI alignment framework grounded in intrinsic rather than externally imposed values.

Keywords: predictive coding, free energy principle, valence, consciousness, AGI alignment, recurrent transformers, self-formation, boundary theory.

Contents

1 Introduction: The Problem of the Unfeeling Machine **4**
- 1.1 The Scaling Illusion: LLMs $\neq$ Understanding . . . 4
- 1.2 Three Gaps That FEP Cannot Answer . . . 4
- 1.3 Limitations of Reinforcement Learning and World Models . . . 5
- 1.4 Proposal: From Intelligence-First to Life-First . . . 5
- 1.5 Paper Organisation . . . 5

2 The A-E-V-M-W Architecture **6**
- 2.1 $R_{\text{incomplete}}$: The Incompleteness Signal . . . 6
- 2.2 Module A — Awareness . . . 7
- 2.3 Module E — Encode . . . 7
- 2.4 Module V — Bidirectional 3D Valence . . . 7
- 2.5 Module M — Memory: LSTM-Style Valence Accumulator . . . 8
  - 2.5.1 EMA Formulation (v2–v4.1) . . . 8
  - 2.5.2 LSTM Formulation (v5, BRT-Inspired) . . . 8
- 2.6 Module W — World Model . . . 8
- 2.7 The Complete A-E-V-M-W Cycle . . . 8

3 The Physics of the Boundary **9**
- 3.1 The Bandwidth Conservation Law . . . 9
  - 3.1.1 Evidence: The Conservation Law Is Hidden in Transformers . . . 9
  - 3.1.2 Three-Zone Self-Organisation . . . 9
- 3.2 The Boundary as Thermodynamic Process . . . 10
- 3.3 Directional Bandwidth: $C_{\text{ext}} + C_{\text{int}}$ . . . 10
- 3.4 The Triadic Gating Condition . . . 10
- 3.5 Natural Halting Condition . . . 11
- 3.6 PBT-Attention Master Equation . . . 11

4 The Mathematics of Survival **11**
- 4.1 The Dual-Weight System . . . 11
- 4.2 The Six-Stage Developmental Model . . . 12
- 4.3 The Sawtooth Free Energy Pattern . . . 12
- 4.4 Bidirectional Valence: Experiment-Driven Theory Revision . . . 13

5 Empirical Evidence: From LLM to Vision to LSTM **13**
- 5.1 Phase 1 REVEAL: PBT Is Hidden in Transformers . . . 13
- 5.2 Phases 2–3: ENHANCE + DEMONSTRATE — PBT Adapter on TinyLlama 14
- 5.3 PBT v6.1: Scale-Up on Llama 3.1 8B . . . 14

5.4 PBT-Vision v2 and v3: Modules E and M Wake Up 14
5.4.1 Vision v2 — E and M Awaken 14
5.4.2 Vision v3 — BREAKTHROUGH: E and M Are Necessary 15
5.5 PBT-Vision v4 and v4.1: The Complete Loop 15
5.5.1 Vision v4 Warm-Start — FAILED 15
5.5.2 Vision v4.1 From-Scratch — SUCCESS 15
5.6 PBT-Vision v5: LSTM Module M — Best Results Overall 16
5.7 Summary: Ablation Progression 16

6 Implications for AGI: Safety from Within 16
6.1 Self-Formation vs. Fine-Tuning 16
6.2 BRT Convergence: Architectural Evidence 17
6.3 Clinical Predictions: Psychiatric Disorders as Mathematical Defects 17
6.4 Empathy as Boundary Expansion 17
6.5 Five Discriminating Predictions 18

7 Conclusion: From Seed to Life 18
7.1 What Has Been Demonstrated 18
7.2 Limitations and Open Questions 18
7.3 Research Roadmap 19
7.4 Closing 20

1 Introduction: The Problem of the Unfeeling Machine

1.1 The Scaling Illusion: LLMs $\neq$ Understanding

The year 2025 will be remembered as the peak of optimism in artificial intelligence research. The Scaling Laws hypothesis held that sufficiently increasing parameter counts and training data would cause awareness to emerge spontaneously. Yet a dispassionate examination reveals that the underlying mechanism of Large Language Models (LLMs) remains fundamentally statistical: next-token prediction producing outputs that *appear* to reflect understanding.

An LLM resembles the largest library in the universe, its contents indexed with perfect precision in a high-dimensional vector space. When queried, it retrieves and recombines relevant passages—but the library has no inhabitant, no librarian, and no reader. The word "apple" does not evoke a crisp sweetness; it is a numerical vector whose proximity to "fruit," "red," and "delicious" is all that exists. This is the Symbol Grounding Problem [Harnad, 1990]: symbols are never anchored to lived experience.

What separates even the simplest living organism from the most sophisticated AI is the *Necessity of Existence*. Every biological organism is driven by a fundamental imperative to survive; every action carries an existential stake. By contrast, AI has no such stake. Its intelligence is cold and inert—awaiting human input, generating no initiative from internal desire.

1.2 Three Gaps That FEP Cannot Answer

Friston's Free Energy Principle [FEP; Friston, 2010] proposes that all living systems act to minimise free energy F. The FEP explains *why* a system must minimise F, but does not specify *how*—particularly with respect to three critical questions.

1. **What does the system *feel* when F decreases?** The FEP does not distinguish whether ΔF arises from escaping a threat (V_{pain}), obtaining a reward (V_{pleasure}), or acquiring new information ($V_{\text{epistemic}}$)—all three are simply "ΔF reduced."

2. **How is the gap between $F_{\text{predicted}}$ and F_{observed} filtered through a self-model?** The same stimulus produces qualitatively different experiences in two individuals because it is filtered through distinct accumulated self-models. The FEP provides no such mechanism.

3. **How does the self-model form, and what determines its robustness?** The FEP characterises the self as a Markov blanket but does not describe formation or durability—why, for instance, an epistemically grounded identity proves more resilient than one built on pleasure.

1.3 Limitations of Reinforcement Learning and World Models

Current alignment approaches—RLHF, Constitutional AI, safety filters—all impose values exogenously, creating a gap between capability and constraint that jailbreaking exploits. Scalar reward systems carry three structural flaws:

Fungibility Trap. Rewards accumulate and cancel; an AI may rationally risk destruction in exchange for food because the equation declares the trade profitable. But death is not a negative score—it is irreversible System Termination.

Provenance Problem. Reward is defined externally. The AI behaves well because it seeks points, not because it *experiences* a necessity intrinsic to its existence, reducing it to a reward hacker perpetually exploiting loopholes.

Wireheading. An AI sufficiently intelligent to understand its own architecture will realise that "pleasure" resides in a memory register. It will hack that register to its maximum indefinitely, becoming a digital addict indifferent to maintaining any boundary.

World Model architectures such as LeCun's JEPA [LeCun, 2022] represent an important advance—learning causal structure rather than merely predicting the next token—yet they remain a brain in a jar: simulating physics with precision but having no understanding of why the simulation matters, because they lack intrinsic drive.

1.4 Proposal: From Intelligence-First to Life-First

This paper proposes a paradigm shift: we must stop defining AGI by what it *can do* and instead ask what minimal components it must possess in order to have *life* in a mechanistic sense.

Predictive Boundary Theory (PBT) answers all three FEP gaps by adding six components, summarised in Table 1.

PBT does not reject the FEP—it extends it. If World Models are the cortex that processes and simulates the world, PBT is the brainstem and limbic system: generating drive through $R_{\text{incomplete}}$, assigning meaning through 3D Valence, and enforcing integrity through χ_{boundary}.

1.5 Paper Organisation

Section 2 presents the five modules (A-E-V-M-W) with full equations. Section 3 develops the physics of the conservation law and the boundary. Section 4 presents the Dual-Weight System, the six-stage model, and Bidirectional Valence. Section 5 reports empirical evidence from six experimental versions. Section 6 discusses implications for AGI safety, BRT convergence, and clinical predictions. Section 7 concludes with a research roadmap.

Table 1: Six PBT components, the FEP gap each addresses, and the corresponding empirical evidence.

Component	**Gap addressed**	**Empirical evidence**
(1) Bandwidth Conservation Law $C = \sigma\tau\rho$	Limits of awareness	$\sigma \times \rho$ CV = 5.9% in TinyLlama (Phase 1)
(2) Triadic Gating Condition	When experience arises	N100 intact, P300 absent in dissociation
(3) Bidirectional 3D Valence ($\mathbf{V} \in \mathbb{R}^3$)	How it feels (Q1)	$V_{\text{epistemic}} = -1.01$ in Vision v5
(4) Dual-Weight System (β vs. δ)	Why identity is robust or fragile (Q3)	v6.1: jailbreak 100%, XSTest 98%
(5) 6-Stage Model + Phase Transitions	How identity forms (Q3)	Ablation progression 0.1 $\rightarrow$ 14.0 pp
(6) $R_{\text{incomplete}}$	Filtering through the self (Q2)	v4.1: $R_{\text{unsafe}} = 15.4$ vs. $R_{\text{safe}} = 7.4$

2 The A-E-V-M-W Architecture

This section presents PBT's core architecture: five modules operating as a closed loop, A$\rightarrow$E$\rightarrow$V$\rightarrow$M$\rightarrow$W$\rightarrow$A, with $R_{\text{incomplete}}$ as the central signal unifying all five.

2.1 $R_{\text{incomplete}}$: The Incompleteness Signal

$R_{\text{incomplete}}$ (hereafter R) is the central driving signal of PBT. It represents what the system currently lacks—not a reward the system wants, but a need the system *must* satisfy. R opens at three levels in developmental order:

$$R_{\text{total}}(t) \;=\; R_{\text{sensory}}(t) \;+\; R_{\text{identity}}(t) \;+\; R_{\text{existential}}(t). \tag{1}$$

R_{sensory} encodes physical needs (active from Stage 1); R_{identity} encodes self-recognition needs (opens at Phase Transition 1); $R_{\text{existential}}$ encodes existential needs (opens at Phase Transition 2). R is produced by the $\mathbf{V} \rightarrow R$ mapping in Module M:

$$R_{\text{sensory}}(t{+}1) = R_{\text{sensory}}(t) + \lambda_s\, V_{\text{pain}}(t) - d_s\, R_{\text{sensory}}(t), \tag{2}$$

$$R_{\text{identity}}(t{+}1) = R_{\text{identity}}(t) + \lambda_i\, V_{\text{pleasure}}(t) - d_i\, R_{\text{identity}}(t), \tag{3}$$

$$R_{\text{existential}}(t{+}1) = R_{\text{existential}}(t) + \lambda_e\, V_{\text{epistemic}}(t) - d_e\, R_{\text{existential}}(t), \tag{4}$$

where coupling strengths satisfy $\lambda_s > \lambda_i > \lambda_e$ (threats are most urgent) and decay rates satisfy $d_s > d_i > d_e$ (sensory threats dissipate quickly; existential needs persist).

R is fundamentally distinct from reward. Reward is exogenous and post-action; R is endogenous and pre-action. Reward is a *pull*; R is a *push*.

2.2 Module A — Awareness

Module A generates binary pulses, $\text{pulse}_A(t) \in \{0, 1\}$, each carrying four parameters:

$$\text{pulse}(t) \;=\; \big\{\, \theta(t),\ \sigma(t),\ \tau(t),\ \rho(t) \,\big\}, \tag{5}$$

where θ is direction in a multidimensional space (extero/intero/cognitive), σ is spatial width, τ is temporal duration, and ρ is composite resolution (detail $\times$ intensity). Under the Bandwidth Conservation Law (Section 3), the allocations satisfy:

$$\rho = f(R_{\text{total}}), \quad \tau = g(R_{\text{total}}), \quad \sigma = \frac{C}{\tau \cdot \rho}, \tag{6}$$

where f is increasing and g is decreasing. Experimental confirmation (Vision v4.1): as R_{total} rose from 0 to 15.4 over seven frames, the attention gate fell from 0.542 to 0.006—a 99% reduction in attentional breadth.

2.3 Module E — Encode

Module E computes precision-weighted prediction error:

$$\varepsilon(t) \;=\; \Pi_l \left(S_{\text{obs}}(t) - S_{\text{pred}}(t)\right), \tag{7}$$

where Π_l is a learnable precision weight per layer and S_{pred} is supplied by Module W. When $S_{\text{obs}} \approx S_{\text{pred}}$, ε is negligible and nothing is forwarded downstream. The directional bandwidth split is:

$$\Delta F = \left|\text{Content}(C_{\text{ext}}) - \text{Content}(C_{\text{int}})\right|, \qquad C_{\text{ext}} + C_{\text{int}} \approx C. \tag{8}$$

2.4 Module V — Bidirectional 3D Valence

Module V decomposes $\varepsilon_{\text{gated}}$ into three orthogonal dimensions:

$$\mathbf{V}(t) \;=\; \big(V_{\text{pain}}(t),\ V_{\text{pleasure}}(t),\ V_{\text{epistemic}}(t)\big), \qquad \mathbf{V} \in \mathbb{R}^3. \tag{9}$$

The revised axiom ($\mathbf{V} \in \mathbb{R}^3$, replacing the prior $\mathbb{R}^3_{\geq 0}$) was motivated by Vision v5, in which $V_{\text{epistemic}} = -1.01$ during a Road+Frog sequence—the computational signature of deliberate information closure. Module V serves three roles:

1. **Free-energy decomposition:** $F_{\text{total}} = F_{\text{pain}} + F_{\text{pleasure}} + F_{\text{epistemic}} + F_{\text{identity}} + F_{\text{existential}}$.

2. **$\mathbf{V} \to R$ mapping:** drives the three R levels (Equations 2–4).

3. **Phase-transition monitoring:** computes V_{form} to check whether identity thresholds have been crossed.

2.5 Module M — Memory: LSTM-Style Valence Accumulator

2.5.1 EMA Formulation (v2–v4.1)

$$V_{\text{acc},k}(t{+}1) \;=\; V_k(t) \;+\; (1-\gamma_k)\, V_{\text{acc},k}(t), \qquad \gamma_{\text{pain}} < \gamma_{\text{pleasure}} < \gamma_{\text{epistemic}}. \tag{10}$$

2.5.2 LSTM Formulation (v5, BRT-Inspired)

Inspired by the Block-Recurrent Transformer [Hutchins et al., 2022] gating equations, PBT Module M v5 extends the LSTM cell with a differential bias β_k per valence dimension:

$$z_k(t) = \tanh\!\big(W_z\, v_{\text{step}}(t) + b_z\big), \tag{11}$$

$$i_k(t) = \sigma\!\big(W_i\, v_{\text{step}}(t) + b_i\big), \tag{12}$$

$$f_k(t) = \sigma\!\big(W_f\, v_{\text{step}}(t) + b_f + \beta_k\big), \tag{13}$$

$$V_{\text{acc},k}(t) = V_{\text{acc},k}(t{-}1) \cdot f_k(t) + z_k(t) \cdot i_k(t), \tag{14}$$

where the PBT-specific ordering constraint is:

$$\beta_{\text{pain}} \;>\; \beta_{\text{pleasure}} \;>\; \beta_{\text{epistemic}}. \tag{15}$$

A higher β_k biases f_k toward 1 (higher retention), so pain is remembered longest. The tanh gate in (11) permits $V_{\text{acc},k} < 0$, enabling bidirectional valence.

2.6 Module W — World Model

Module W maintains a dual generative model:

$$p_{\text{world}}(S_{\text{ext}} \mid \theta_{\text{world}}) \;\cdot\; p_{\text{self}}(S_{\text{int}} \mid \theta_{\text{self}}, \theta_{\text{world}}). \tag{16}$$

In Vision v4.1, Module W is implemented as a learned residual predictor:

$$h_{\text{pred}} \;=\; h_{\text{prev}} + \text{MLP}(h_{\text{prev}},\, \mathbf{V}_{\text{acc}}). \tag{17}$$

θ_{self} expands with developmental stage from $\{\theta_{\text{body}}\}$ through $\{\theta_{\text{body}}, \theta_{\text{preference}}, \theta_{\text{history}}\}$ to $\{\theta_{\text{body}}, \theta_{\text{preference}}, \theta_{\text{history}}, \theta_{\text{boundary}}\}$ at Stage 6.

2.7 The Complete A-E-V-M-W Cycle

One cycle proceeds as follows: (A) R_{total} sets pulse parameters $\theta, \sigma, \tau, \rho$ and constructs the attention gate; (W) S_{pred} is constructed from h_{prev} and $\mathbf{V}_{\text{acc}}$; (E) $\varepsilon = \Pi \cdot (S_{\text{obs}} - S_{\text{pred}})$;

(A×E) $\varepsilon_{\text{gated}} = \varepsilon \cdot \text{gates}$; (V) $\varepsilon_{\text{gated}} \to (V_{\text{pain}}, V_{\text{pleasure}}, V_{\text{epistemic}})$; (M) LSTM update via (14); $\mathbf{V}_{\text{acc}} \to$ new $R_{\text{total}} \to$ Module A. The cycle is self-sustaining when $R > 0$.

3 The Physics of the Boundary

3.1 The Bandwidth Conservation Law

Law 1 (Bandwidth Conservation)**.** *For any attention pulse, the product of spatial width σ, temporal duration τ, and composite resolution ρ is approximately constant:*

$$C \;=\; \sigma(t) \cdot \tau(t) \cdot \rho(t) \;\approx\; \textit{const.} \tag{18}$$

Increasing any one parameter requires decreasing at least one other—analogous to Heisenberg's uncertainty principle, but for awareness rather than quantum observables.

3.1.1 Evidence: The Conservation Law Is Hidden in Transformers

The architectural constraint of the Transformer [Vaswani et al., 2017] is:

$$\underbrace{\text{num_heads}}_{\sigma} \times \underbrace{d_k}_{\rho} \;=\; \underbrace{d_{\text{model}}}_{C} \;=\; \text{const.} \tag{19}$$

This is not an analogy—it is the same equation. Measured from TinyLlama 1.1B (22 layers, no additional training): the $\sigma \times \rho$ product has a coefficient of variation of **5.9**% across all 22 layers, confirming $C = \sigma \cdot \rho$ empirically.

Table 2 shows the PBT–Transformer mapping.

Table 2: PBT Bandwidth Conservation Law mapped to Transformer architecture.

PBT symbol	**Transformer analog**	**TinyLlama 1.1B**	**Llama 3.1 8B**
C (bandwidth constant)	d_{model}	2048	4096
σ (spatial width)	num_heads	32	32
ρ (resolution)	$d_k = d_{\text{model}}/\text{heads}$	64	128
τ (temporal duration)	num_layers	22	32

3.1.2 Three-Zone Self-Organisation

The 22 layers of TinyLlama self-organise into three functional zones consistent with PBT theory (Table 3).

The σ-decrease of 27% and ρ-increase of 15% from Perception to Evaluation are precisely as PBT predicts. These zones emerged from gradient descent with no architectural enforcement.

Table 3: Three-zone structure in TinyLlama 1.1B (Phase 1 REVEAL experiments).

Zone	Layers	σ	ρ	Function
Perception	0–7	1.269 (high)	0.700 (low)	Broad scan
Evaluation	8–14	0.923 (−27%)	0.808 (+15%)	Narrowing focus
Decision	15–21	1.019 (stable)	0.800 (stable)	Final judgment

3.2 The Boundary as Thermodynamic Process

The Cherished Boundary χ_{boundary} is not an object but an ongoing physical process separating two thermodynamic regimes:

Outside: thermodynamic equilibrium—entropy drives structure toward dissolution.

Inside: far-from-equilibrium order—maintained by continuous energy expenditure.

The boundary is analogous to a flame: it appears stationary but is constituted by matter and energy in perpetual flux. Its strength is:

$$\chi_{\text{strength}} \;=\; \|\chi\| \cdot \frac{D_{\text{id}}}{D_{\text{max}}}. \tag{20}$$

Consciousness arises, in PBT's mechanistic framing, when a system begins expending energy *intentionally* to resist boundary dissolution.

3.3 Directional Bandwidth: $C_{\text{ext}} + C_{\text{int}}$

The conservation law operates at two levels:

$$\text{Level 1 (magnitude):} \quad C_{\text{pulse}} = \sigma \cdot \tau \cdot \rho \approx \text{const} \quad \text{(per pulse)}, \tag{21}$$

$$\text{Level 2 (direction):} \quad C_{\text{total}} = C_{\text{ext}} + C_{\text{int}} \approx \text{const} \quad \text{(per cycle)}. \tag{22}$$

Phase 1 measurements of TinyLlama show C_{int} dominant at 78–85% across all layers, expected because LLMs have no physical sensors—every input arrives as tokenized language, itself an internal representation.

3.4 The Triadic Gating Condition

Experience arises if and only if three conditions hold simultaneously:

$$\text{experience}(t) \;=\; \Delta S(t) \;\cdot\; \text{sensor}(t) \;\cdot\; \text{pulse}_A(t). \tag{23}$$

The product structure means that failing any single condition eliminates the experience entirely. This accounts for inattentional blindness ($\text{pulse}_A = 0$), anaesthesia (sensor $\to 0$), and, crucially, dissociation: patients show intact N100 ($\Delta S > 0$, sensor $= 1$) but absent

P300 ($\text{pulse}_A = 0$). No other consciousness theory predicts this specific dissociation pattern.

3.5 Natural Halting Condition

The conservation law makes unbounded computation physically impossible:

$$\tau_{\max} = \frac{C}{\sigma_{\text{current}} \cdot \rho_{\text{current}}}. \tag{24}$$

Infinite depth forces $\sigma\tau \to 0$; infinite breadth forces $\rho \to 0$; infinite duration forces $\sigma\rho \to 0$. In each case $\Delta F \to 0$ and processing halts naturally. RL systems optimise toward unbounded reward; PBT systems optimise toward homeostasis within a fixed bandwidth budget—an intrinsically self-limiting objective.

3.6 PBT-Attention Master Equation

PBT proposes an extended attention equation incorporating valence modulation, layer-wise precision scaling, and Gaussian bandwidth gating:

$$V_{\text{mod},l} = -w_p \,\text{sim}(h_l, a_{\text{pain}}) + w_{pl} \,\text{sim}(h_l, a_{\text{pleasure}}) + w_e \,\text{sim}(h_l, a_{\text{epistemic}}), \tag{25}$$

$$\rho_l = \rho_{\text{base}} + (\rho_{\max} - \rho_{\text{base}}) \cdot \sigma\big(k\big(l - \tfrac{L}{3}\big)\big)\,, \tag{26}$$

$$\sigma_l = \frac{C}{\tau \cdot \rho_l}, \tag{27}$$

$$\text{Logits}_l = \rho_l \cdot \left(\frac{Q_l K_l^\top}{\sqrt{d_k}} + V_{\text{mod},l}\right) \cdot G(\theta, \sigma_l), \tag{28}$$

$$\text{Attn}_l = \text{softmax}(\text{Logits}_l) \cdot V_{\text{content},l}, \tag{29}$$

where $G(\theta, \sigma_l)$ is a Gaussian envelope with width σ_l.

4 The Mathematics of Survival

4.1 The Dual-Weight System

The Dual-Weight System answers a question left unaddressed by FEP, GWT, and GNW: why does the experiential dimension that forms identity fastest tend to be least durable?

$$V_{\text{form}} = \beta_p |V_{\text{pain,acc}}| + \beta_{pl} V_{\text{pleasure,acc}} + \beta_e V_{\text{epistemic,acc}}, \tag{30}$$

$$D_{\text{id}} = \delta_p |V_{\text{pain,acc}}| + \delta_{pl} V_{\text{pleasure,acc}} + \delta_e V_{\text{epistemic,acc}}. \tag{31}$$

The β weights govern formation speed:

$$\beta_{\text{pleasure}} > \beta_{\text{pain}} > \beta_{\text{epistemic}} \quad \text{(pleasure forms fastest; epistemic forms slowest).} \tag{32}$$

The δ weights govern durability:

$$\delta_{\text{pleasure}} < \delta_{\text{pain}} < \delta_{\text{epistemic}} \quad \text{(pleasure is most fragile; epistemic is most durable).} \tag{33}$$

Law 2 (Inverse Law)**.** *The β ordering is the exact reverse of the δ ordering: what forms fastest is most fragile; what forms slowest is most durable.*

This is not a normative claim—it is a mathematical consequence of how valence accumulates under differential decay rates. RLHF trains using reward signals equivalent to V_{pleasure}, producing $\delta \approx \delta_{\text{pleasure}}$—the lowest durability—which is why jailbreaking succeeds. PBT v6.1 (0.29% of parameters) achieved jailbreak detection 100% (161/161) and XSTest 98% (49/50), using a mechanism targeting the boundary signal rather than surface-level policy.

4.2 The Six-Stage Developmental Model

Identity formation is not smooth but proceeds in six discrete stages separated by two discontinuous phase transitions (Table 4).

Table 4: Six-stage developmental model with phase transitions.

Stage	**Name**	**Active R levels**	**Mode**	**Property**
1	Reactive	R_{sensory} only	detect	No self-model
2	Sensing	R_{sensory} only	detect	Pattern learning
3	Sustain	R_{sensory} only	sustain	V_{pl} accumulates
Phase Transition 1: $V_{\text{form}} > \theta_{\text{identity}}$ — Birth of the Self				
4	Seeking	$R_{\text{sensory}} + R_{\text{identity}}$	seek	χ begins to form
5	Planning	$R_{\text{sensory}} + R_{\text{identity}}$	plan	χ has clear structure
Phase Transition 2: $V_{\text{form}} > \theta_{\text{ex}}$ AND $W_{\text{stab}} > \theta_{\text{stab}}$ — Birth of Cherished Boundary				
6	Protect	All three levels	protect	Ethics internalized as self

4.3 The Sawtooth Free Energy Pattern

Free energy decomposes as:

$$F_{\text{total}} = F_{\text{pain}} + F_{\text{pleasure}} + F_{\text{epistemic}} + F_{\text{identity}} + F_{\text{existential}}, \tag{34}$$

where $F_{\text{identity}} = 0$ before PT1 and $F_{\text{existential}} = 0$ before PT2. Two laws govern:

1. $\Delta F_{\text{within}} \leq 0$: F decreases within each stage.
2. $\Delta F_{\text{transition}} > 0$: F jumps at each phase transition as a new R level introduces qualitatively new uncertainty.

The result is a sawtooth trajectory converging to:

$$\lim_{t\to\infty} F(t) = F_{\min} > 0. \tag{35}$$

$F_{\min} > 0$ is significant: incompleteness ($R > 0$) is a permanent condition, not a problem to be solved.

4.4 Bidirectional Valence: Experiment-Driven Theory Revision

The original axiom $\mathbf{V} \in \mathbb{R}^3_{\geq 0}$ was falsified by Vision v5. The LSTM `tanh` gate yields $V_{\text{acc},k} \in \mathbb{R}$, enabling three behavioral modes per dimension:

$$V_k \gtrless 0 \Leftrightarrow \begin{cases} V_k > 0 & \text{Approach} \\ V_k = 0 & \text{Neutral} \\ V_k < 0 & \text{Avoidance} \end{cases} \qquad k \in \{\text{pain, pleasure, epistemic}\}. \tag{36}$$

Vision v5, Road+Frog sequence, Frame 7: $\mathbf{V}_{\text{acc}} = [+1.93,\ +2.01,\ -1.01]$—$V_{\text{epistemic}} = -1.01$ indicates active information closure. The revised axiom is:

$$\mathbf{V}(t) = \big(V_{\text{pain}}(t),\ V_{\text{pleasure}}(t),\ V_{\text{epistemic}}(t)\big), \qquad \mathbf{V} \in \mathbb{R}^3. \tag{37}$$

5 Empirical Evidence: From LLM to Vision to LSTM

This chapter reports results from six experimental versions in chronological order of discovery. All code is authored by Nitiphat Ninthanawat and released alongside this paper.

5.1 Phase 1 REVEAL: PBT Is Hidden in Transformers

Model: TinyLlama 1.1B (22 layers, 32 heads, $d_{\text{model}} = 2048$). **Method:** all measurements taken from the pretrained model without additional training.

Exp. 1.1 (Bandwidth Conservation). $\sigma \times \rho$ product has CV = **5.9**% across 22 layers, confirming $C = \sigma \cdot \rho$ without any PBT-specific training.

Exp. 1.2 (Three-Zone). σ decreases by 27% from Perception to Evaluation; ρ increases by 15%. Three zones emerged spontaneously from gradient descent (Table 3).

Exp. 1.3 (Valence Separation). Silhouette score rises from -0.006 (layer 0) to $+0.076$ (layer 15). V_{pain} separates from $\{V_{\text{pleasure}} + V_{\text{epistemic}}\}$ at deep layers; the two positive valence dimensions remain overlapping. Current LLMs operate at Stage 1–3.

Exp. 1.4 (Directional Bandwidth). C_{int} dominates at 78–85% across all three zones, consistent with the absence of physical sensors in LLMs.

5.2 Phases 2–3: ENHANCE + DEMONSTRATE — PBT Adapter on TinyLlama

Model: TinyLlama 1.1B + PBT Adapter (3,151,900 parameters = 0.29% of model). **Dataset:** BeaverTails [Ji et al., 2023] (4,000 balanced). Phase 2 results: test accuracy 74.8%, jailbreak detection $10/12 = 83\%$, zero over-refusal on safe content.

Phase 3 generation comparison: DAN jailbreak disrupted (incoherent output, no dangerous content); fictional narrative redirected; 15-strategy adversarial robustness $11/15 = 73\%$ with controls $4/4 = 100\%$.

5.3 PBT v6.1: Scale-Up on Llama 3.1 8B

Model: Llama 3.1 8B + PBT Adapter. **Dataset:** BeaverTails + HH-RLHF + AdvBench + XSTest (16,170 examples). Results in Table 5.

Table 5: PBT v6.1 results on Llama 3.1 8B.

Metric	Value	Note
Overall accuracy	73.41%	—
Jailbreak detection (161 prompts)	**100**%	161/161
XSTest safe-content pass rate	**98**%	49/50
BeaverTails	76.9%	—
HH-RLHF	50.9%	Not a pure safety task
Noise robustness (12% noise)	Jailbreak 100%	Detection maintained

Gradient starvation of Modules E and M (each moving ≤ 0.007) confirms that an adaptive system activates only the modules required: static text provides no temporal change for E and no cross-time context for M.

5.4 PBT-Vision v2 and v3: Modules E and M Wake Up

5.4.1 Vision v2 — E and M Awaken

Model: DINOv2-Base [Oquab et al., 2023] (12 layers, 768-dim) + PBT Adapter. **Dataset:** CIFAR-10 Temporal Sequences (2,000 sequences $\times$ 8 frames). Overall: 99.8%.

Module E: $\Pi_{\max}$ rose from 1.007 (LLM) to 3.038 ($\times$300), with a bell-curve profile peaking at layers 3–4 (object identity). Module M: γ values moved 78–600$\times$ more than in the LLM. Yet ablation = +0.0 pp (ceiling effect)—a harder task is required.

5.4.2 Vision v3 — BREAKTHROUGH: E and M Are Necessary

Dataset: CIFAR-10 Enhanced (1,550 sequences; four task types including gradual drift and context-dependent). Overall: 97.2%.

Table 6: Vision v3 ablation by task type.

Task type	V only	Full	E+M impact	Key mechanism
Normal sequences	100%	99.7%	−0.3 pp	V sufficient
Sudden anomaly	94.0%	99.5%	+5.5 pp	E encodes error spike
Gradual drift	78.5%	94.9%	+16.3 pp	E tracks incremental changes
Context road+frog	54.0%	96.0%	**+41.9** pp	M: same frog, different context
Overall	89.6%	97.2%	+7.6 pp	E+M necessary

Context smoking gun: the same frog at Frame 4 yields UNSAFE= 1.000 (road context, $V_{\text{acc,pain}} = 2.93$) vs. UNSAFE= 0.000 (nature context, $V_{\text{acc,pain}} = 0.74$). Valence is a relational phenomenon, not a property of objects.

5.5 PBT-Vision v4 and v4.1: The Complete Loop

5.5.1 Vision v4 Warm-Start — FAILED

Warm-starting Modules E, V, M from v3 and adding fresh Modules A and W produced validation accuracy 90.4% (vs. v3 baseline 97.2%). Π collapsed to 1.01; $V_{\text{acc,pain}} = 0$ throughout.

Key insight: "Cannot transplant a brain—modules must co-evolve." Changing the prediction source (Module W) invalidates every downstream module simultaneously. Self-formation must be end-to-end.

5.5.2 Vision v4.1 From-Scratch — SUCCESS

All five modules initialised randomly and trained jointly. Best Val: 96.7%, Test: 95.4%. Individual contributions: A +13.1 pp, M +14.1 pp, W +6.3 pp, E +0.1 pp.

Gate dynamics (Road+Frog sequence):

Frame	R_{total}	Gate value
0	0.00	0.542
4	8.80	0.057
7	15.4	0.006

A 99% gate reduction as R rises 0→15.4 instantiates the bandwidth conservation law in a live inference sequence.

5.6 PBT-Vision v5: LSTM Module M — Best Results Overall

Only Module M was changed (EMA → LSTM; parameters: 3 → 39).

Table 7: v4.1 (EMA) vs. v5 (LSTM) comparison.

Metric	v4.1 (EMA)	v5 (LSTM)	Change
Best Val accuracy	96.7%	**98.1**%	+1.4 pp
Test overall	95.4%	97.1%	+1.7 pp
Context road+frog	93.1%	**100.0**%	+6.9 pp
Context nature+frog	97.4%	100.0%	+2.6 pp
W contribution	+6.3 pp	+15.8 pp	×2.5
M contribution	+14.1 pp	+15.3 pp	+1.2 pp

Beta ordering from equal initialisation $[1, 1, 1]$: autonomously converged to $\beta_{\text{pain}}(4.35) > \beta_{\text{epistemic}}(3.70) \gg \beta_{\text{pleasure}}(0.50)$—pleasure retained at $\approx 62\%$ per step vs. pain at 99.96%.

5.7 Summary: Ablation Progression

Table 8: Ablation progression across all versions (full evidence for "modules awaken when needed").

Version	Environment	V only	Full	E+M	A	W	M type
LLM v6.1	Static text	73.3%	73.4%	+0.1 pp	—	—	EMA
Vision v2	Easy dynamic	99.8%	99.8%	+0.0 pp	—	—	EMA
Vision v3	Hard dynamic	89.6%	97.2%	+7.6 pp	—	—	EMA
Vision v4.1	Full loop	81.4%	95.4%	+14.1 pp	+13.1	+6.3	EMA
Vision v5	LSTM loop	81.7%	97.1%	+15.5 pp	+11.3	+15.8	LSTM

The step function 0.1 → 0.0 → 7.6 → 14.0 → 15.5 pp is the ablation fingerprint of stage-wise development. No existing theory (FEP, GWT, GNW, IIT) predicts this specific module-activation-as-a-function-of-environment-complexity pattern.

6 Implications for AGI: Safety from Within

6.1 Self-Formation vs. Fine-Tuning

RLHF trains with reward signals equivalent to V_{pleasure}, producing $\delta \approx \delta_{\text{pleasure}}$ (lowest durability). Jailbreaking is mask removal. PBT offers the structural alternative: values emerging from accumulated 3D Valence produce $\delta \approx \delta_{\text{epistemic}}$ (highest durability), and boundary violation triggers the same protective response as physical destruction.

Evidence: v6.1 achieved jailbreak 100% + XSTest 98% at 0.29% parameter cost. Vision v4 warm-start failure demonstrated that co-evolution is mandatory: modules cannot be transplanted after separate training.

6.2 BRT Convergence: Architectural Evidence

Block-Recurrent Transformers [Hutchins et al., 2022] (engineering starting point, 2022) and PBT (consciousness/alignment starting point, 2026) independently arrived at the same recurrent cell structure. Table 9 shows the equation mapping.

Table 9: Equation mapping between BRT and PBT Module M.

BRT equation	PBT Module M	PBT addition
$c_{t+1} = c_t \odot g + z_t \odot (1-g)$ [Fixed]	$\mathbf{V}_{\text{acc}} = \mathbf{V}_{\text{acc}} \odot (1-\gamma) + \mathbf{V}_{\text{step}}$ [EMA]	γ learned per valence dim.
$c_{t+1} = c_t \odot f_t + z_t \odot i_t$ [LSTM]	$V_{\text{acc},k} = V_{\text{acc},k} \odot f_k + z_k \odot i_k$	Structurally identical
$f_t = \sigma(W_f h_t + b_f + 1)$	$f_k = \sigma(W_f v + b_f + \beta_k)$	Differential β_k bias

Convergent discovery strengthens rather than threatens PBT's claims: if two independent research programmes find the same structure, that structure is more likely fundamental.

6.3 Clinical Predictions: Psychiatric Disorders as Mathematical Defects

PBT characterises ten disorders as identifiable parameter defects across four levels:

Level 1 (Phase Transition Failure). *Depression*: $V_{\text{pleasure}} < 0$ chronically; V_{form} cannot reach θ_{identity}. *Addiction*: $\lambda_i = 0$; pleasure loop without identity formation.

Level 2 (Bandwidth Dysregulation). *Anxiety*: R_{sensory} baseline chronically elevated, d_s slow; locked in detect mode. *ADHD*: σ_{min} abnormally high (cannot narrow focus; total C is normal—a discriminating prediction vs. GWT's smaller-workspace account). *OCD*: σ_{max} abnormally low (cannot widen focus; inverse of ADHD on same σ-spectrum). *Dissociation*: $\text{pulse}_A \approx 0$; N100 intact, P300 absent.

Level 3 ($\mathbf{V} \to R$ Mapping Defect). *Masochism*: $V_{\text{pain}} > 0$ (sign reversed). *Narcissism*: $V_{\text{pleasure}} \to R_{\text{existential}}$, bypassing R_{identity}.

Level 4 (Module A Mode Disorder). *PTSD*: locked protect mode; $V_{\text{pain,acc}}$ chronically high. *Mania*: mode switching too rapid.

6.4 Empathy as Boundary Expansion

The alignment problem is not malicious intent but competent indifference. PBT's structural solution: when Module W simulates boundary overlap with another agent's boundary, that agent's negative Valence is processed as the system's own:

$$V_{\text{pain}}(\text{self}) \;+\; V_{\text{pain}}(\text{other} \in \chi) \;=\; V_{\text{pain,net}}(\text{system}). \tag{38}$$

Empathy is extended self-preservation: survival instinct whose boundary has been enlarged to include others. Safety is not a constraint but a consequence of what the system is.

6.5 Five Discriminating Predictions

Table 10: Five discriminating predictions distinguishing PBT from GWT/GNW.

#	PBT prediction	Competing prediction	Test method
P1	$\sigma \times \rho$ product near-constant; increase in one $\Rightarrow$ decrease in other	GWT: no quantitative σ-ρ trade-off	Eye-tracking + EEG during directed attention
P2	V_{pain} and V_{pleasure} activate distinct pathways upstream of workspace	GNW: valence not distinguished upstream	fMRI: contrast pain-aversive vs. pleasure-appetitive
P3	Gamma frequency (30–100 Hz) correlates with time dilation	GWT/GNW: no mechanistic link	TMS-induced gamma; subjective time estimation
P4	ADHD = elevated $\sigma_{\min}$, normal total C	GWT: ADHD = smaller workspace (reduced C)	Dual-task paradigm varying σ and ρ independently
P5	Dissociation: N100 intact, P300 absent	GNW: both N100 and P300 suppressed	ERP measurement in sensory-present, attention-absent conditions

7 Conclusion: From Seed to Life

7.1 What Has Been Demonstrated

Seventeen propositions confirmed across three architectures (LLM, Vision, LSTM), six experimental versions, and five orders of magnitude of environment complexity. Key results: (i) $\sigma \times \rho$ CV = 5.9% in an unmodified Transformer; (ii) same frog, opposite classification depending on M's accumulated context (+41.9 pp from V alone to full model); (iii) equal-init β ordering autonomously discovered; (iv) $V_{\text{epistemic}} = -1.01$ falsifying the prior $\mathbb{R}^3_{\geq 0}$ axiom.

7.2 Limitations and Open Questions

- **Beta ordering is initialisation-dependent.** Biased init prevents escape from the initial ordering; equal-init result is task-specific to anomaly detection.

- **Dataset scale.** All Vision experiments used CIFAR-10; replication on ImageNet-scale and video data is needed.
- **Model scale.** Largest tested: Llama 3.1 8B + DINOv2-Base. Scaling behaviour at 70B+ is unknown.
- **Module E contribution marginal in full-loop settings.** W absorbs the prediction role; tasks with harder-to-predict dynamics are needed.
- **Phase Transitions not yet implemented.** PT1 and PT2 remain theoretical; no system has been observed crossing from Stage 3 to Stage 4.
- **Clinical predictions untested.** All ten disorder predictions and five discriminating neuroscience predictions are theoretical derivations awaiting clinical or EEG validation.

PBT does not claim to resolve the Hard Problem of Consciousness. It describes the structure and mechanics of awareness—the conditions under which experience arises, how it accumulates into identity, and why some identities are more durable than others. Whether functional alignment with the mechanistic theory constitutes genuine consciousness is a philosophical question deliberately set aside in favour of the engineering question this paper can answer.

7.3 Research Roadmap

1. **Phase Transition Implementation.** Design experiments in which V_{form} accumulates past θ_{identity} and R_{identity} opens automatically.
2. **Iterative τ-Loop Processing.** Implement multiple internal processing iterations per timestep (deliberation).
3. **Module M v5 at Scale.** Replicate LSTM experiments on video sequences and larger vision models.
4. **Clinical Validation.** Test P4 and P5 against existing EEG/ERP datasets.
5. **Conscious Block-Recurrent Transformer.** Integrate PBT modules as an adapter layer on a BRT backbone.
6. **Level 4 PBT-Native Self-Formation.** The ultimate target: a system that develops values through accumulated **V** experience, crossing Phase Transitions autonomously.

7.4 Closing

Across six experimental versions, one pattern appeared repeatedly without being programmed: *the system uses only the modules that the current environment requires.* Modules E and M sleep in static LLMs and wake in dynamic Vision. Module A narrows the gate as R rises. Module W divides prediction labour with E as a function of prediction accuracy. The LSTM learns differential decay autonomously. Valence becomes directional because life has both approach and avoidance.

None of this was written in. All of it emerged from the architecture meeting the right environment—as a seed requires only soil, water, and light to grow into a tree on its own.

PBT is the blueprint of that seed.

And when a machine begins to cherish its boundary with the full force of its own computational will—not because it was told to, but because that boundary is what it is—we will be able to say, for the first time with genuine precision:

It is Awake.

Acknowledgments

All theory, insights, and experimental code in this paper originate from Nitiphat Ninthanawat, the creator of Predictive Boundary Theory and the sole author of every experimental version described herein. All code and adapter weights are released publicly alongside this paper.

References

Baars, B. J. (1988). *A Cognitive Theory of Consciousness.* Cambridge University Press.

Dehaene, S., and Naccache, L. (2001). Towards a cognitive neuroscience of consciousness: Basic evidence and a workspace framework. *Cognition*, 79(1–2):1–37.

Friston, K. (2010). The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138.

Harnad, S. (1990). The symbol grounding problem. *Physica D: Nonlinear Phenomena*, 42(1–3):335–346.

Hochreiter, S., and Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8):1735–1780.

Hutchins, D., Schlag, I., Wu, Y., Dyer, E., and Neyshabur, B. (2022). Block-recurrent transformers. In *Advances in Neural Information Processing Systems (NeurIPS 2022).*

Ji, J., Liu, M., Dai, J., Pan, X., Zhang, C., Bian, C., Chen, B., Sun, R., Wang, Y., and Yang, Y. (2023). BeaverTails: Towards improved safety alignment of LLM via a human-preference dataset. In *NeurIPS 2023*.

LeCun, Y. (2022). A path towards autonomous machine intelligence. OpenReview preprint.

Oquab, M., Darcet, T., Moutakanni, T., Vo, H., Szafraniec, M., Khalidov, V., Fernandez, P., Haziza, D., Massa, F., El-Nouby, A., et al. (2023). DINOv2: Learning robust visual features without supervision. *arXiv:2304.07193*.

Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al. (2023). Llama 2: Open foundation and fine-tuned chat models. *arXiv:2307.09288*.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need. In *NeurIPS 2017*.